

Shahd Mohamed Abd El Sabour

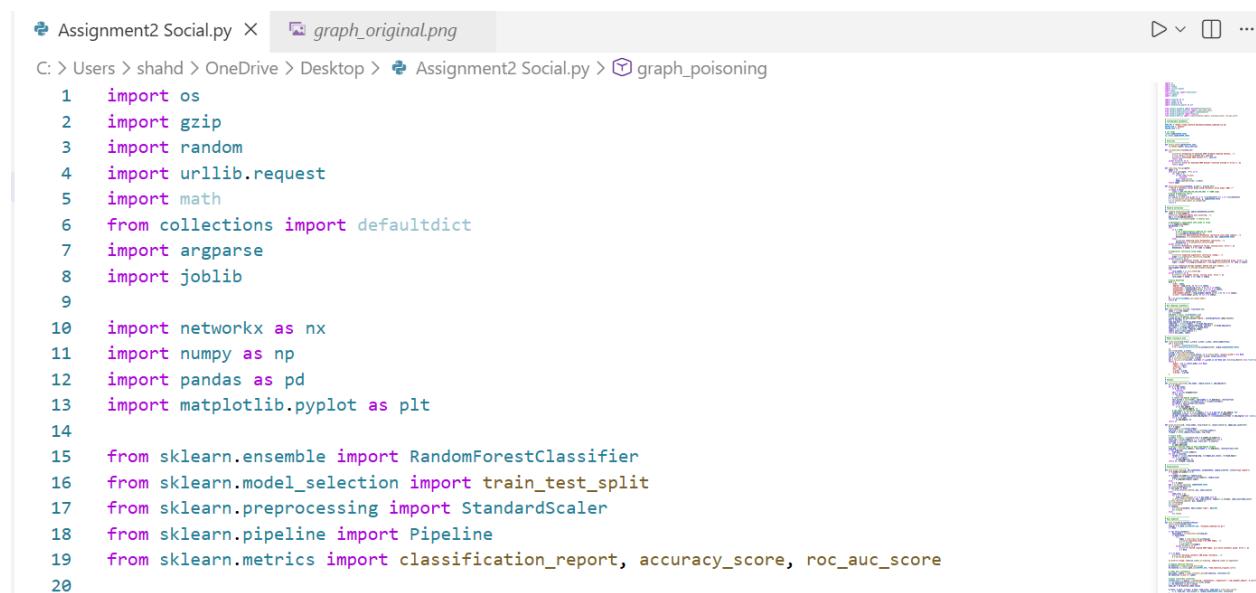
2205027

Assignment 2 Social Networking

Overview :

This assignment builds a social graph (tries to download the SNAP Facebook network, otherwise creates a synthetic SBM graph), computes per-node graph features, synthetically labels a small fraction of nodes as “bots”, trains a Random Forest classifier to detect bots, then simulates two adversarial attacks:

- **Structural evasion:** bots change their connectivity (remove some edges to high-degree neighbors and add edges to low-degree nodes) to try to evade detection.
- **Graph poisoning:** the training data is corrupted by flipping labels and injecting fake nodes with many edges; the model is retrained on this poisoned data and evaluated on the clean test set.



The screenshot shows a code editor window with the following details:

- Title Bar:** Assignment2_Social.py X graph_original.png
- File Path:** C: > Users > shahd > OneDrive > Desktop > Assignment2_Social.py > graph_poisoning
- Code Content (Lines 1-20):**

```
1 import os
2 import gzip
3 import random
4 import urllib.request
5 import math
6 from collections import defaultdict
7 import argparse
8 import joblib
9
10 import networkx as nx
11 import numpy as np
12 import pandas as pd
13 import matplotlib.pyplot as plt
14
15 from sklearn.ensemble import RandomForestClassifier
16 from sklearn.model_selection import train_test_split
17 from sklearn.preprocessing import StandardScaler
18 from sklearn.pipeline import Pipeline
19 from sklearn.metrics import classification_report, accuracy_score, roc_auc_score
20
```
- Right Panel:** A large network graph visualization showing a complex social network structure with many nodes and edges, color-coded according to node features.

- Standard libs for files, randomness, downloading.
- networkx for graph operations.
- pandas/numpy for data handling.
- sklearn for modeling and metrics.
- joblib to save models.

```

16 from sklearn.model_selection import train_test_split
17 from sklearn.preprocessing import StandardScaler
18 from sklearn.pipeline import Pipeline
19 from sklearn.metrics import classification_report, accuracy_score, roc_auc_score
20
21 # -----
22 # Configurable parameters
23 # -----
24 SNAP_URL = "https://snap.stanford.edu/data/facebook_combined.txt.gz"
25 OUTPUT_DIR = "outputs"
26 RANDOM_SEED = 42
--
```

RANDOM_SEED used to make behavior reproducible (where supported). The script also sets Python and numpy seeds.

```

31
32 # -----
33 # Utilities
34 # -----
35 def ensure_outdir(path=OUTPUT_DIR):
36     os.makedirs(path, exist_ok=True)
37
38 def try_download_snap(save_to):
39     try:
40         print("[*] Attempting to download SNAP facebook_combined dataset...")
41         urllib.request.urlretrieve(SNAP_URL, save_to)
42         print("[+] Downloaded SNAP dataset to:", save_to)
43         return True
44     except Exception as e:
45         print("[!] Could not download SNAP dataset (internet blocked or error):", e)
46         return False
47
48 def load_snap_from_gz(path):
49     edges = []
50     with gzip.open(path, "rt") as f:
51         for line in f:
52             if not line.strip():
53                 continue
54             a,b = line.split()
55             edges.append((int(a), int(b)))
56     return edges
--
```

`ensure_outdir(path=OUTPUT_DIR)`

Creates the output directory if it doesn't exist.

`try_download_snap(save_to)`

Attempts to download the SNAP facebook dataset to `save_to`. Returns True on success, False on any exception (e.g., no internet).

`load_snap_from_gz(path)`

Reads the gzipped edge list line by line and returns a list of (a,b) integer edges. Used if the SNAP download succeeds.

```
54     |     |     a,b = line.split()
55     |     |     edges.append((int(a), int(b)))
56     return edges
57
58 def build_sbm_graph(sizes=None, p_in=0.1, p_out=0.005):
59     """Build a synthetic social graph using stochastic block model (SBM)."""
60     if sizes is None:
61         sizes = [400,350,300,250,200,200,300] # ~2000 nodes
62     # Build probability matrix
63     nblocks = len(sizes)
64     p = [[p_in if i==j else p_out for j in range(nblocks)] for i in range(nblocks)]
65     G = nx.stochastic_block_model(sizes, p, seed=RANDOM_SEED)
66     G = nx.convert_node_labels_to_integers(G)
67     return G
68
```



`build_sbm_graph(sizes=None, p_in=0.1, p_out=0.005)`

- Builds a Stochastic Block Model (SBM) with `sizes` blocks (default block sizes sum to ~2000 nodes).
- Constructs a probability matrix `p` where intra-block edge prob = `p_in` and inter-block prob = `p_out`.
- Calls `nx.stochastic_block_model` and re-labels nodes to consecutive integers.
- Returns a `networkx.Graph`.

Use-case: fallback if SNAP download or load fails.

Assignment2 Social.py X graph_original.png

C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > compute_graph_metrics

```
68
69  # -----
70  # Feature extraction
71  #
72 def compute_graph_metrics(G, approx_betweenness_k=200):
73     nodes = list(G.nodes())
74     print("[*] Computing degree and clustering...")
75     deg = dict(G.degree(nodes))
76     clustering = nx.clustering(G)  # returns dict
77
78     # Betweenness: approximate when graph is large
79     n = G.number_of_nodes()
80     betweenness = {}
81     try:
82         if n > 1000:
83             # use k-approximation sampling for speed
84             k = min(approx_betweenness_k, n)
85             print(f"[*] Approximating betweenness centrality with k={k} samples...")
86             betweenness = nx.betweenness_centrality(G, k=k, seed=RANDOM_SEED)
87         else:
88             print("[*] Computing exact betweenness centrality...")
89             betweenness = nx.betweenness_centrality(G)
90     except Exception as e:
91         print("![!] Betweenness computation failed, setting zeros. Error:", e)
92         betweenness = {node: 0.0 for node in nodes}
93
94     # Eigenvector centrality using numpy
95     try:
96         print("[*] Computing eigenvector centrality (numpy)...")
97         eigen = nx.eigenvector_centrality_numpy(G)
```

In 73, Col 26 Spaces: 4 UTE-8 CRLE [1] Python ⚙ Signed out 3.13.9 (Microsoft Store) ⌂

Assignment2 Social.py X graph_original.png

C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > compute_graph_metrics

```
72 def compute_graph_metrics(G, approx_betweenness_k=200):
73     print("[*] Computing eigenvector centrality (numpy)...")
74     eigen = nx.eigenvector_centrality_numpy(G)
75     except Exception as e:
76         print("![!] Eigenvector failed, falling back to degree-normalized proxy. Error:", e)
77         eigen = {node: float(deg.get(node, 0)) / (max(deg.values())+1e-6) for node in nodes}
78
79     print("[*] Computing average neighbor degree and core numbers...")
80     avg_neighbor_degree = nx.average_neighbor_degree(G)
81     try:
82         core_number = nx.core_number(G)
83     except Exception as e:
84         print("![!] core_number failed, filling zeros. Error:", e)
85         core_number = {node: 0 for node in nodes}
86
87     # Build DataFrame
88     data = [
89         'node': nodes,
90         'degree': [deg.get(n, 0) for n in nodes],
91         'clustering': [clustering.get(n, 0.0) for n in nodes],
92         'betweenness': [betweenness.get(n, 0.0) for n in nodes],
93         'eigenvector': [eigen.get(n, 0.0) for n in nodes],
94         'avg_neighbor_degree': [avg_neighbor_degree.get(n, 0.0) for n in nodes],
95         'k_core': [core_number.get(n, 0) for n in nodes],
96     ]
97     df = pd.DataFrame(data).set_index('node')
98     return df
99
100    # -----
101    # Bot labeling (synthetic)
```

In 119, Col 6 Spaces: 4 UTE-8 CRLE [1] Python ⚙ Signed out 3.13.9 (Microsoft Store) ⌂

`compute_graph_metrics(G, approx_betweenness_k=200)`

For graph G, computes per-node features and returns a pandas.DataFrame indexed by node id. Features:

- `degree`: node degree (number of neighbors).
- `clustering`: clustering coefficient (triangles / possible triangles).
- `betweenness`: betweenness centrality. If $n > 1000$, uses k-sample approximate betweenness via `nx.betweenness_centrality(G, k=k, seed=RANDOM_SEED)` for speed; otherwise computes exact.
 - The code catches exceptions and falls back to zeros on failure.
- `eigenvector`: eigenvector centrality using `nx.eigenvector_centrality_numpy`. On failure falls back to a proxy: degree normalized by max degree.
- `avg_neighbor_degree`: average neighbor degree (from NetworkX).
- `k_core`: core number (`nx.core_number`). On failure fills zeros.

Important performance notes:

- Exact betweenness and eigenvector centralities can be expensive on large graphs; the script attempts approximate betweenness when $n > 1000$.

Return: DataFrame with one row per node and columns listed above.

```
121     return df
122
123 # -----
124 # Bot labeling (synthetic)
125 #
126 def label_synthetic_bots(df, fraction=0.05):
127     nodes = list(df.index)
128     n = len(nodes)
129     num_bots = max(1, int(fraction * n))
130     # pick half top-degree, half random
131     sorted_by_deg = df.sort_values('degree', ascending=False).index.tolist()
132     half = num_bots // 2
133     high_deg_bots = sorted_by_deg[:half]
134     remaining = list(set(nodes) - set(high_deg_bots))
135     random_bots = random.sample(remaining, num_bots - len(high_deg_bots))
136     bot_nodes = set(high_deg_bots + random_bots)
137     labels = pd.Series(0, index=df.index)
138     labels.loc[list(bot_nodes)] = 1
139     return bot_nodes, labels
140
```

`label_synthetic_bots(df, fraction=0.05)`

- `fraction` fraction of nodes are labeled as bots (default 5%).
- Strategy: choose half of bot nodes as the highest-degree nodes, and the other half randomly from the rest.
- Returns `bot_nodes` (a set) and `labels` (a pandas.Series indexed by nodes with 1 for bot, 0 otherwise).

This creates a plausible synthetic scenario where some bots are high-degree (influencer-like) and others look random.

```
139     return bot_nodes, labels
140
141 # -----
142 # Model training & eval
143 # -----
144 def train_evaluate(X_train, y_train, X_test, y_test, return_model=True):
145     clf = Pipeline([
146         ('scaler', StandardScaler()),
147         ('rf', RandomForestClassifier(n_estimators=200, random_state=RANDOM_SEED))
148     ])
149     clf.fit(X_train, y_train)
150     y_pred = clf.predict(X_test)
151     y_proba = clf.predict_proba(X_test)[:,1] if hasattr(clf, "predict_proba") else None
152     report = classification_report(y_test, y_pred, output_dict=True)
153     acc = accuracy_score(y_test, y_pred)
154     auc = roc_auc_score(y_test, y_proba) if y_proba is not None and len(set(y_test))>1 else float('na
155     return {
156         'model': clf if return_model else None,
157         'report': report,
158         'accuracy': acc,
159         'auc': auc,
160         'y_pred': y_pred,
161         'y_proba': y_proba
162     }
```

`train_evaluate(X_train, y_train, X_test, y_test, return_model=True)`

- Builds a Pipeline with `StandardScaler` + `RandomForestClassifier(n_estimators=200, random_state=RANDOM_SEED)`.
- Trains on training data, predicts on test data.
- Computes `classification_report` (as dict), accuracy, and AUC (if possible).
- Returns dictionary containing model, report, accuracy, auc, `y_pred`, `y_proba`.

Note: if all test labels are the same, ROC AUC is not computable and the code uses nan.

```
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > structural_evasion
162     }
163
164     # -----
165     # Attacks
166     #
167     def structural_evasion(G, bot_nodes, remove_frac=0.5, add_degree=5):
168         G2 = G.copy()
169         for b in bot_nodes:
170             if b not in G2:
171                 continue
172             nbrs = list(G2.neighbors(b))
173             if not nbrs:
174                 continue
175             # remove top-degree neighbors
176             nbrs_sorted = sorted(nbrs, key=lambda x: G2.degree[x], reverse=True)
177             num_remove = max(1, int(remove_frac * len(nbrs_sorted)))
178             to_remove = nbrs_sorted[:num_remove]
179             for r in to_remove:
180                 if G2.has_edge(b, r):
181                     G2.remove_edge(b, r)
182             # add edges to low-degree nodes
183             candidates = [n for n in G2.nodes() if n != b and not G2.has_edge(b, n)]
184             candidates_sorted = sorted(candidates, key=lambda x: G2.degree[x])
185             to_add = candidates_sorted[:add_degree] if len(candidates_sorted) >= add_degree else random.sample(candidates, add_degree)
186             for a in to_add:
187                 G2.add_edge(b, a)
188
189     return G2
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
```

```
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > graph_poisoning
167     def structural_evasion(G, bot_nodes, remove_frac=0.5, add_degree=5):
168         return G2
169
170     def graph_poisoning(G, train_nodes, flip_frac=0.15, inject_frac=0.02, edges_per_inject=20):
171         G2 = G.copy()
172         train_nodes = list(train_nodes)
173         num_flip = max(1, int(flip_frac * len(train_nodes)))
174         flipped = random.sample(train_nodes, num_flip)
175
176         # Inject nodes
177         n_inject = max(1, int(inject_frac * G.number_of_nodes()))
178         start_idx = max(G2.nodes()) + 1 if len(G2.nodes()) > 0 else 0
179         injected = list(range(start_idx, start_idx + n_inject))
180         for node in injected:
181             G2.add_node(node)
182         # Connect injected nodes to many high-degree targets
183         high_deg = sorted(G2.nodes(), key=lambda x: G2.degree[x], reverse=True)[:200]
184         if not high_deg:
185             high_deg = list(G2.nodes())
186         for inj in injected:
187             targets = random.sample(high_deg, min(edges_per_inject, len(high_deg)))
188             for t in targets:
189                 G2.add_edge(inj, t)
190
191     return G2, flipped, injected
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
```

```
structural_evasion(G, bot_nodes, remove_frac=0.5,  
add_degree=5)
```

- For each bot b:
 - Get its neighbors and sort them by neighbor degree descending.
 - Remove the top `remove_frac` fraction of these neighbors (i.e., disconnect from high-degree neighbors).
 - Add `add_degree` new edges from b to low-degree nodes it wasn't connected to (candidates sorted by degree ascending). If there are not enough low-degree nodes, it picks a random subset.
- Returns a new graph `G2` with the structural changes.
- Goal: make bot nodes less obviously high-degree/hub-linked and more blended with low-degree nodes.

```
graph_poisoning(G, train_nodes, flip_frac=0.15,  
inject_frac=0.02, edges_per_inject=20)
```

- Copies the graph to `G2`.
- Chooses `num_flip = flip_frac * len(train_nodes)` training nodes and marks them as flipped (this list is returned but not immediately applied to labels here — the caller flips labels in the training set).
- Injects `n_inject = inject_frac * number_of_nodes` new nodes with consecutive integer ids beyond current max node id.
- Connects each injected node to many high-degree targets (top 200 high-degree nodes) with up to `edges_per_inject` edges each.
- Returns (`G2, flipped, injected`).

Purpose: emulate poisoning the training graph by injecting noisy nodes and prepare to flip labels for some training nodes.

C: > Users > shabd > OneDrive > Desktop > Assignment2 Social.py > plot_graph_sample

```
211
212 # -----
213 # Visualization
214 #
215 def plot_graph_sample(G, bot_nodes=None, outpath=None, sample_size=500, title="Graph sample"):
216     if G.number_of_nodes() == 0:
217         return
218     if G.number_of_nodes() > sample_size:
219         sample = random.sample(list(G.nodes()), sample_size)
220         H = G.subgraph(sample).copy()
221     else:
222         H = G.copy()
223     pos = nx.spring_layout(H, seed=RANDOM_SEED)
224     plt.figure(figsize=(8,8))
225     if bot_nodes is None:
226         nx.draw_networkx_nodes(H, pos, node_size=20)
227     else:
228         node_color = []
229         for n in H.nodes():
230             node_color.append(1 if n in bot_nodes else 0)
231         nx.draw_networkx_nodes(H, pos, node_size=20, cmap=plt.cm.viridis, node_color=node_color)
232     nx.draw_networkx_edges(H, pos, alpha=0.3)
233     plt.title(title)
234     plt.axis('off')
235     if outpath:
236         plt.savefig(outpath, bbox_inches='tight', dpi=200)
237         plt.close()
238     else:
239         plt.show()
240
```

```
plot_graph_sample(G, bot_nodes=None, outpath=None,
sample_size=500, title="Graph sample")
```

- If the graph has more nodes than `sample_size`, sample `sample_size` node IDs and work with the induced subgraph to keep drawing fast.
- Uses `nx.spring_layout` (seeded) to get positions and draws nodes and edges.
- When `bot_nodes` is given, nodes are colored by whether they're in the `bot_nodes` set (passed values 0/1 to `node_color`).
- Saves to `outpath` if provided, otherwise shows plot.

```
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > main_flow
242 # Main pipeline
243 # -----
244 def main_flow(force_synthetic=False):
245     ensure_outdir(OUTPUT_DIR)
246     snap_gz = os.path.join(OUTPUT_DIR, "facebook_combined.txt.gz")
247     G = None
248
249     if not force_synthetic:
250         downloaded = try_download_snap(snap_gz)
251         if downloaded:
252             try:
253                 edges = load_snap_from_gz(snap_gz)
254                 print("[*] Building graph from SNAP edges...")
255                 G = nx.Graph()
256                 G.add_edges_from(edges)
257             except Exception as e:
258                 print("![!] Failed loading SNAP edges, will build synthetic graph. Error:", e)
259                 G = None
260
261     if G is None:
262         print("[*] Building synthetic SBM graph (fallback)...")
263         G = build_sbm_graph()
264
265     print(f"[+] Graph: nodes={G.number_of_nodes()}, edges={G.number_of_edges()}")
266
267     # Compute baseline features
268     df_features = compute_graph_metrics(G)
269     df_features.to_csv(os.path.join(OUTPUT_DIR, "node_features_original.csv"))
270
271     # Label bots (synthetic)
```

```
Assignment2 Social.py X graph_original.png
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > main_flow
244 def main_flow(force_synthetic=False):
245
246     # Label bots (synthetic)
247     bot_nodes, labels = label_synthetic_bots(df_features, fraction=0.05)
248     df_features['is_bot'] = labels
249
250     # Split train/test (stratify)
251     FEATURE_COLS = ['degree', 'clustering', 'betweenness', 'eigenvector', 'avg_neighbor_degree', 'k_core']
252     X = df_features[FEATURE_COLS].fillna(0).values
253     y = df_features['is_bot'].values
254     node_ids = df_features.index.values
255
256     X_train, X_test, y_train, y_test, node_train, node_test = train_test_split(
257         X, y, node_ids, test_size=0.3, random_state=RANDOM_SEED, stratify=y
258     )
259
260     # Baseline model
261     print("[*] Training baseline classifier...")
262     baseline = train_evaluate(X_train, y_train, X_test, y_test, return_model=True)
263     joblib.dump(baseline['model'], os.path.join(OUTPUT_DIR, "baseline_model.joblib"))
264
265     # Save baseline report
266     with open(os.path.join(OUTPUT_DIR, "baseline_results.txt"), "w") as f:
267         f.write(classification_report(y_test, baseline['y_pred']))
268         f.write(f"\nAccuracy: {baseline['accuracy']}\nAUC: {baseline['auc']}\\n")
269
270     print("[+] Baseline accuracy:", baseline['accuracy'], "AUC:", baseline['auc'])
271
272     # Visualize original graph (sample)
273     plot_graph_sample(G, bot_nodes=bot_nodes, outpath=os.path.join(OUTPUT_DIR, "graph_original.png"),
274
```

```
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > main_flow
244 def main_flow(force_synthetic=False):
245     joblib.dump(baseline['model'], os.path.join(OUTPUT_DIR, "baseline_model.joblib"))
246
247     # Save baseline report
248     with open(os.path.join(OUTPUT_DIR, "baseline_results.txt"), "w") as f:
249         f.write(classification_report(y_test, baseline['y_pred']))
250         f.write(f"\nAccuracy: {baseline['accuracy']}\nAUC: {baseline['auc']}\n")
251
252     print("[+] Baseline accuracy:", baseline['accuracy'], "AUC:", baseline['auc'])
253
254     # Visualize original graph (sample)
255     plot_graph_sample(G, bot_nodes=bot_nodes, outpath=os.path.join(OUTPUT_DIR, "graph_original.png"),
256                       node_size=100, edge_width=1)
257
258     # -----
259     # Structural evasion
260     # -----
261
262     attacked_subset = random.sample(list(bot_nodes), max(1, int(0.3*len(bot_nodes)))) # attack 30% of nodes
263     G_struct = structural_evasion(G, attacked_subset, remove_frac=0.5, add_degree=5)
264     df_struct = compute_graph_metrics(G_struct)
265
266     # Ensure ordering same as original nodes for comparison
267     df_struct = df_struct.reindex(df_features.index).fillna(0)
268     df_struct.to_csv(os.path.join(OUTPUT_DIR, "node_features_structural.csv"))
269
270     # Evaluate baseline model on features computed after structural evasion (no retraining)
271     X_struct_test = df_struct.loc[node_test, FEATURE_COLS].fillna(0).values
272     struct_eval = {}
273
274     # reuse baseline model to predict
275     model_baseline = baseline['model']
276     y_struct_pred = model_baseline.predict(X_struct_test)
277
278     try:
279         y_struct_proba = model_baseline.predict_proba(X_struct_test)[:,1]
280     except Exception:
281         y_struct_proba = None
282
283     struct_report = classification_report(y_test, y_struct_pred, output_dict=False)
284     struct_acc = accuracy_score(y_test, y_struct_pred)
285     struct_auc = roc_auc_score(y_test, y_struct_proba) if y_struct_proba is not None and len(set(y_struct_proba)) > 1 else None
286
287     with open(os.path.join(OUTPUT_DIR, "structural_attack_results.txt"), "w") as f:
288         f.write(struct_report)
289         f.write(f"\nAccuracy: {struct_acc}\nAUC: {struct_auc}\n")
290
291     print("[+] After structural evasion -> accuracy:", struct_acc, "AUC:", struct_auc)
292     plot_graph_sample(G_struct, bot_nodes=bot_nodes, outpath=os.path.join(OUTPUT_DIR, "graph_structural.png"),
293                       node_size=100, edge_width=1)
294
295     # -----
296     # Graph poisoning
297     # -----
298
299     # We'll poison training set: flip some labels and inject nodes, then retrain
300     G_poisoned, flipped_nodes, injected_nodes = graph_poisoning(G, node_train, flip_frac=0.15, inject=True)
301     df_poison = compute_graph_metrics(G_poisoned)
302     df_poison = df_poison.reindex(df_features.index).fillna(0)
303     df_poison.to_csv(os.path.join(OUTPUT_DIR, "node_features_poison.csv"))
304
305     # Prepare poisoned training labels: flip labels for 'flipped_nodes' that are in train
```

```
Ln 311, Col 26  Spaces: 4  UTF-8  CRLF  {} Python  Signed out  3.13.9 (Microsoft Store)  🔍
C: / Users / Shahd / OneDrive / Desktop / Assignment2 Social.py > main_flow > mapped_nodes
244 def main_flow(force_synthetic=False):
245     struct_eval = {}
246
247     # reuse baseline model to predict
248     model_baseline = baseline['model']
249     y_struct_pred = model_baseline.predict(X_struct_test)
250
251     try:
252         y_struct_proba = model_baseline.predict_proba(X_struct_test)[:,1]
253     except Exception:
254         y_struct_proba = None
255
256     struct_report = classification_report(y_test, y_struct_pred, output_dict=False)
257     struct_acc = accuracy_score(y_test, y_struct_pred)
258     struct_auc = roc_auc_score(y_test, y_struct_proba) if y_struct_proba is not None and len(set(y_struct_proba)) > 1 else None
259
260     with open(os.path.join(OUTPUT_DIR, "structural_attack_results.txt"), "w") as f:
261         f.write(struct_report)
262         f.write(f"\nAccuracy: {struct_acc}\nAUC: {struct_auc}\n")
263
264     print("[+] After structural evasion -> accuracy:", struct_acc, "AUC:", struct_auc)
265     plot_graph_sample(G_struct, bot_nodes=bot_nodes, outpath=os.path.join(OUTPUT_DIR, "graph_structural.png"),
266                       node_size=100, edge_width=1)
267
268     # -----
269     # Graph poisoning
270     # -----
271
272     # We'll poison training set: flip some labels and inject nodes, then retrain
273     G_poisoned, flipped_nodes, injected_nodes = graph_poisoning(G, node_train, flip_frac=0.15, inject=True)
274     df_poison = compute_graph_metrics(G_poisoned)
275     df_poison = df_poison.reindex(df_features.index).fillna(0)
276     df_poison.to_csv(os.path.join(OUTPUT_DIR, "node_features_poison.csv"))
277
278     # Prepare poisoned training labels: flip labels for 'flipped_nodes' that are in train
```

```

Assignment2 Social.py X graph_original.png
C: > Users > shahd > OneDrive > Desktop > Assignment2 Social.py > main_flow
244 def main_flow(force_synthetic=False):
339
340     # Prepare poisoned training labels: flip labels for 'flipped_nodes' that are in train
341     # Find positions within node_train
342     y_train_poison = y_train.copy()
343     node_train_list = list(node_train)
344     for n in flipped_nodes:
345         if n in node_train_list:
346             idx = node_train_list.index(n)
347             y_train_poison[idx] = 1 - y_train_poison[idx] # flip
348
349     # Build poisoned training X using new features computed on poisoned graph
350     train_positions = [list(df_poison.index).index(n) for n in node_train] # mapping node id -> row
351     X_train_poison = df_poison.iloc[train_positions][FEATURE_COLS].fillna(0).values
352
353     # Retrain classifier on poisoned data
354     clf_poisoned = Pipeline([
355         ('scaler', StandardScaler()),
356         ('rf', RandomForestClassifier(n_estimators=200, random_state=RANDOM_SEED))
357     ])
358     print("[*] Training classifier on poisoned training data...")
359     clf_poisoned.fit(X_train_poison, y_train_poison)
360
361     # Evaluate poisoned-trained model on clean test features (from original df_features)
362     X_test_clean = df_features.loc[node_test, FEATURE_COLS].fillna(0).values
363     y_poisoned_pred = clf_poisoned.predict(X_test_clean)
364     try:
365         y_poisoned_proba = clf_poisoned.predict_proba(X_test_clean)[:,1]
366     except Exception:
367         y_poisoned_proba = None

```

- 1. Prepare outputs:** call `ensure_outdir(OUTPUT_DIR)`.
- 2. Download / load graph:**
 - a. If `force_synthetic` is `False`, try to download the SNAP file to `outputs/facebook_combined.txt.gz`.
 - b. If download succeeds, load edges using `load_snap_from_gz` and create `G = nx.Graph()` with `G.add_edges_from(edges)`.
 - c. If any step fails, or if `force_synthetic=True`, fall back to `G = build_sbm_graph()` (the synthetic SBM).
- 3. Print graph size:** nodes and edges.
- 4. Compute baseline features:**
 - a. `df_features = compute_graph_metrics(G)` and save to `node_features_original.csv`.
- 5. Label bots:**
 - a. `bot_nodes, labels = label_synthetic_bots(df_features, fraction=0.05)`.
 - b. Add labels as column `is_bot` in `df_features`.
- 6. Prepare feature matrix and split:**

- a. FEATURE_COLS list:

```
[ 'degree', 'clustering', 'betweenness', 'eigenvector', 'avg_neig
hbor_degree', 'k_core' ].
```
- b. X = feature values, y = is_bot, node_ids = index (node ids).
- c. train_test_split(..., test_size=0.3,
random_state=RANDOM_SEED, stratify=y) returns X_train, X_test,
y_train, y_test, node_train, node_test.
i. NOTE: stratify ensures similar bot fraction in both sets.

7. Train baseline model:

- a. Calls train_evaluate(...) and saves model as
baseline_model.joblib.
- b. Writes a text file baseline_results.txt containing classification report,
accuracy, and AUC.
- c. Prints baseline accuracy and AUC.

8. Visualize original graph:

- a. Calls plot_graph_sample(...,
outpath="outputs/graph_original.png").

9. Structural evasion experiment:

- a. Select attacked_subset — 30% of bot_nodes at random to attack.
- b. G_struct = structural_evasion(G, attacked_subset,
remove_frac=0.5, add_degree=5).
- c. Compute features on G_struct → df_struct and reindex to original index.
- d. Save node_features_structural.csv.
- e. Evaluate the **baseline** model on the evaded features for the test nodes (no
retraining).
 - i. Build X_struct_test from df_struct.loc[node_test,
FEATURE_COLS].
 - ii. y_struct_pred = model_baseline.predict(X_struct_test)
and compute metrics.
- f. Save structural_attack_results.txt and print the results.
- g. Plot sample graph_structural_attack.png.

10. Graph poisoning experiment:

- a. G_poisoned, flipped_nodes, injected_nodes =
graph_poisoning(G, node_train, flip_frac=0.15,
inject_frac=0.02, edges_per_inject=20).
- b. Compute features on G_poisoned → df_poison and reindex.
- c. Save node_features_poison.csv.
- d. **Flip labels** in the training set:

- i. Copy `y_train` to `y_train_poison`.
- ii. For every `n` in `flipped_nodes` that appears in `node_train`, flip the label at that index (`1 -> 0, 0 -> 1`).
- e. Rebuild poisoned training `X_train_poison` using the new features `df_poison` for nodes in `node_train`.
 - i. The code maps nodes in `node_train` to their row positions in `df_poison` using `train_positions = [list(df_poison.index).index(n) for n in node_train]`.
- f. Retrain a fresh RandomForest pipeline (same config) on (`X_train_poison`, `y_train_poison`).
- g. Evaluate this poisoned-trained model on the **clean** test features `X_test_clean = df_features.loc[node_test, FEATURE_COLS]`.
- h. Save `poisoning_attack_results.txt`, dump `poisoned_model.joblib`, and plot `graph_poisoning.png`.
- i. Print results.

11. Save a Markdown report:

- a. `report_summary.md` contains counts and summarized metrics for baseline, structural evasion, and poisoning experiments and lists produced files.

12. Finish: prints that outputs are saved.

```
# -----
# CLI
# -----
def parse_args():
    p = argparse.ArgumentParser(description="Assignment 2: Bot detection + attacks pipeline")
    p.add_argument("--force-synthetic", action="store_true", help="Skip SNAP download and force synthetic data")
    return p.parse_args()

if __name__ == "__main__":
    args = parse_args()
    main_flow(force_synthetic=args.force_synthetic)
```



`parse_args()` uses `argparse` to accept `--force-synthetic`.

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS CODE REFERENCE LOG Python Debug Console + ×

```
rs\shahd\.vscode\extensions\ms-python.debugpy-2025.16.0-win32-x64\bundled\libs\debugpy\launcher' '59837' '--' 'c:\Users\shahd\OneDrive\Desktop\Assignment2 Social.py'
PS C:\Users\shahd\OneDrive\Desktop\outputs> & 'c:\Users\shahd\AppData\Local\Microsoft\WindowsApps\python3.13.exe' 'c:\Users\shahd\.vscode\extensions\ms-python.debugpy-2025.16.0-win32-x64\bundled\libs\debugpy\launcher' '59837' '--' 'c:\Users\shahd\OneDrive\Desktop\Assignment2 Social.py'
[*] Attempting to download SNAP facebook_combined dataset...
[+] Downloaded SNAP dataset to: outputs\facebook_combined.txt.gz
[*] Building graph from SNAP edges...
[+] Graph: nodes=4039, edges=88234
[*] Computing degree and clustering...
[*] Approximating betweenness centrality with k=200 samples...
[*] Computing eigenvector centrality (numpy)...
[*] Computing average neighbor degree and core numbers...
[*] Training baseline classifier...
[+] Baseline accuracy: 0.971947194719472 AUC: 0.7526475694444444
[*] Computing degree and clustering...
[*] Approximating betweenness centrality with k=200 samples...
[*] Computing eigenvector centrality (numpy)...
[*] Computing average neighbor degree and core numbers...
[+] After structural evasion -> accuracy: 0.9570957095709571 AUC: 0.7396122685185185
[*] Computing degree and clustering...
[*] Approximating betweenness centrality with k=200 samples...
[*] Computing eigenvector centrality (numpy)...
[*] Computing average neighbor degree and core numbers...
[*] Training classifier on poisoned training data...
[+] Poisoned training -> accuracy on clean test: 0.9645214521452146 AUC: 0.7565972222222221
[+] All outputs saved to: outputs
PS C:\Users\shahd\OneDrive\Desktop\outputs>
```

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	node	degree	clustering	betweenness	eigenvect	avg_neigh	k_core							
2	0	347	0.041962	0.147377	3.31E-05	18.95965	21							
3	1	17	0.419118	1.56E-06	5.97E-07	48.23529	13							
4	2	10	0.888889	1.55E-07	2.17E-07	49.9	9							
5	3	17	0.632353	2.41E-06	6.56E-07	59.76471	13							
6	4	10	0.866667	0	2.17E-07	42.6	9							
7	5	13	0.333333	7.10E-06	1.18E-06	50.61538	10							
8	6	6	0.933333	0	2.11E-07	63.5	5							
9	7	20	0.431579	8.23E-05	2.57E-05	45.9	12							
10	8	8	0.678571	0	2.13E-07	48.375	5							
11	9	57	0.397243	6.81E-06	2.20E-06	42.40351	21							
12	10	10	0.822222	0	7.64E-07	79.1	10							
13	11	1	0	0	2.04E-07	347	1							
14	12	1	0	0	2.04E-07	347	1							
15	13	31	0.651613	8.21E-07	1.09E-06	54.54839	21							
16	14	15	0.742857	3.83E-07	2.24E-07	38.66667	10							
17	15	1	0	0	2.04E-07	347	1							
18	16	9	0.666667	0	2.55E-07	66.88889	9							
19	17	13	0.730769	3.54E-07	2.21E-07	42.76923	9							
20	18	1	0	0	2.04E-07	347	1							
21	19	16	0.283333	0	2.24E-07	33.125	7							
22	20	15	0.685714	3.61E-07	2.23E-07	37.4	9							
23	21	65	0.349038	0.001332	2.59E-05	42.4	21							
24	22	11	0.472727	0	1.03E-06	53.72727	9							

A	B	C	D	E	F	G	H	I	J	K	L	M
node	degree	clustering	betweenness	eigenvec	avg_neigh	k_core						
0	178	0.010284	0.081509	9.57E-05	5.882022	9						
1	16	0.341667	0.000486	7.92E-07	28.5625	12						
2	10	0.711111	0.002551	6.89E-07	32.2	9						
3	16	0.583333	3.04E-05	2.40E-06	41	12						
4	10	0.866667	0	7.15E-07	25.7	9						
5	12	0.212121	0.000838	3.69E-06	25.08333	9						
6	6	0.933333	0	6.93E-07	35.33333	5						
7	19	0.368421	0.000174	4.52E-05	29.21053	11						
8	8	0.678571	0	7.02E-07	27.25	5						
9	56	0.375325	0.000245	4.01E-06	35.875	20						
10	10	0.622222	9.89E-05	1.79E-06	61.2	9						
11	4	0	0.000436	0.000528	94	4						
12	12	4	0	0.000436	0.000528	94						
13	13	30	0.627586	5.83E-05	1.80E-06	43.66667	20					
14	14	14	0.703297	6.09E-07	2.29E-08	15.71429	9					
15	15	4	0	0.000436	0.000528	94						
16	9	0.444444	8.91E-05	7.62E-07	47.22222	9						
17	12	0.681818	4.02E-06	5.66E-08	16.41667	8						
18	4	0	0.000127	0.000279	73.75	4						
19	19	15	0.180952	0.000146	3.61E-06	11.86667	7					
20	20	14	0.637363	6.24E-07	4.98E-08	14.57143	9					
21	21	64	0.328373	0.007061	4.57E-05	36.54688	20					
22	10	0.355556	5.37E-07	1.65E-06	23.4	8						

A	B	C	D	E	F	G	H	I	J	K	L
node	degree	clustering	betweenness	eigenvec	avg_neigh	k_core					
0	356	0.03988	0.12997	0.000351	18.99719	21					
1	17	0.419118	1.30E-06	2.78E-06	48.76471	13					
2	10	0.888889	0	2.29E-06	50.8	9					
3	17	0.632353	1.17E-06	2.86E-06	60.29412	13					
4	10	0.866667	0	2.28E-06	43.5	9					
5	13	0.333333	3.55E-06	3.30E-06	51.30769	10					
6	6	0.933333	0	2.23E-06	65	5					
7	20	0.431579	5.67E-05	2.79E-05	46.35	12					
8	8	0.678571	0	2.25E-06	49.5	5					
9	57	0.397243	1.30E-06	5.01E-06	42.5614	21					
10	10	0.822222	0	2.87E-06	80	10					
11	1	0	0	2.16E-06	356	1					
12	1	0	0	2.16E-06	356	1					
13	31	0.651613	1.55E-06	3.53E-06	54.83871	21					
14	15	0.742857	0	2.36E-06	39.26667	10					
15	1	0	0	2.16E-06	356	1					
16	9	0.666667	0	2.33E-06	67.88889	9					
17	13	0.730769	0	2.33E-06	43.46154	9					
18	1	0	0	2.16E-06	356	1					
19	16	0.283333	0	2.37E-06	33.6875	7					
20	15	0.685714	0	2.36E-06	38	9					
21	65	0.349038	0.000651	2.89E-05	42.53846	21					
22	11	0.472727	0	3.12E-06	54.54545	9					

baseline_results.txt

File Edit View

	precision	recall	f1-score	support
0	0.97	1.00	0.99	1152
1	0.88	0.50	0.64	60
accuracy			0.97	1212
macro avg	0.93	0.75	0.81	1212
weighted avg	0.97	0.97	0.97	1212

Accuracy: 0.971947194719472
AUC: 0.7526475694444444

structural_attack_results.txt

File Edit View Close tab (Ctrl+W)

	precision	recall	f1-score	support
0	0.96	1.00	0.98	1152
1	1.00	0.13	0.24	60
accuracy			0.96	1212
macro avg	0.98	0.57	0.61	1212
weighted avg	0.96	0.96	0.94	1212

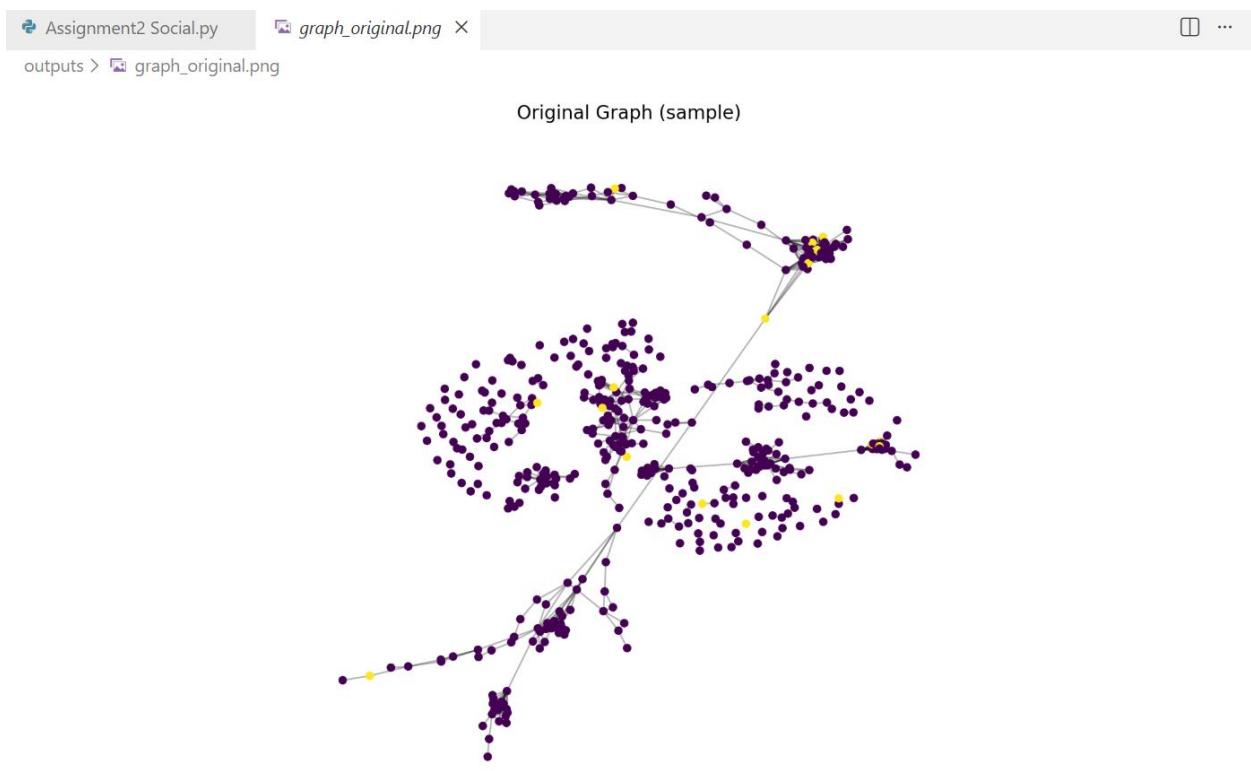
Accuracy: 0.9570957095709571
AUC: 0.7396122685185185

poisoning_attack_results.txt

File Edit View

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1152
1	0.71	0.48	0.57	60
accuracy			0.96	1212
macro avg	0.84	0.74	0.78	1212
weighted avg	0.96	0.96	0.96	1212

Accuracy: 0.9645214521452146
AUC: 0.7565972222222221



Assignment2 Social.py

graph_structural_attack.png X

...

outputs > graph_structural_attack.png

After Structural Evasion (sample)



Whole Image 1280x1314 165.41KB

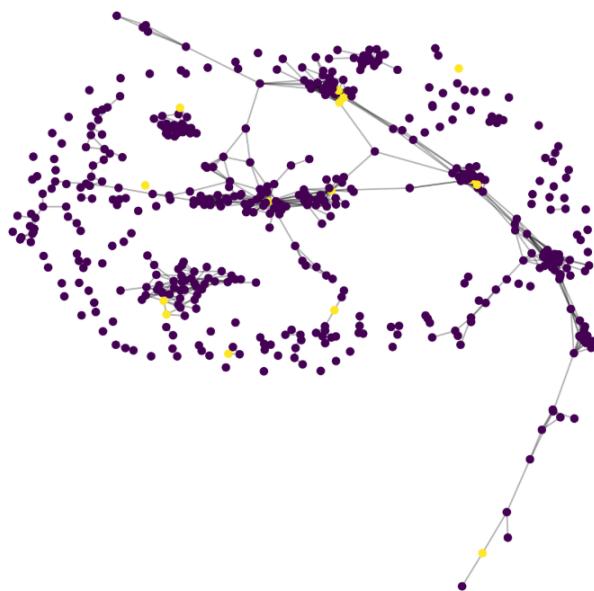
Assignment2 Social.py

graph_poisoning.png X

outputs > graph_poisoning.png

...

After Graph Poisoning (sample)



Whole Image

1280x1314

Signed out

177.44KB

